

Nonblocking I/O

Hyang-Won Lee
Department of Computer Science and Engineering
Konkuk University
leehw@konkuk.ac.kr
<https://sites.google.com/view/leehwko/>

Revisiting Multithreaded Program

- Limitations of naive multithreaded servers
 - one thread for each connection or request \Rightarrow too many threads for server with high volume of requests
 - overhead (thread is not free)
 - spawning threads
 - switching between threads
 - memory
- Solutions
 - thread pool
 - nonblocking I/O

Blocking I/O vs Nonblocking I/O

- `ServerSocket.wait()`, `InputStream.read()`, ...
 - block I/O operations (CPU idled)
- Another thread can grab CPU
 - many threads can help but again threads require resources for management
- Nonblocking I/O (in a single thread)
 - when I/O operation cannot be performed, switch to other operations
 - e.g., server reading from and writing to client
 - when client socket is not ready to be written, try to read from another client

When is Nonblocking I/O useful?

- Example scenario

- a server needs to serve a huge number of long-lived, simultaneous connections, say +10,000
- each client doesn't send much data

- One thread for each connection

- too many active threads

- Thread pool

- can be a bad choice because it takes long until a thread is released to pool (other connections may have to wait for long)

- Nonblocking I/O can be a winner

Example Client

Chargen Client

- RFC 864

- when client connects to server, server keeps sending ASCII characters

```
!"#$%&'()*+,-./0123456789:;<=>?@ABCDEFGHIJKLMNOPQRSTUVWXYZ[\]^_`abcdefgh  
!"#$%&'()*+,-./0123456789:;<=>?@ABCDEFGHIJKLMNOPQRSTUVWXYZ[\]^_`abcdefghi  
!"#$%&'()*+,-./0123456789:;<=>?@ABCDEFGHIJKLMNOPQRSTUVWXYZ[\]^_`abcdefghij  
!"#$%&'()*+,-./0123456789:;<=>?@ABCDEFGHIJKLMNOPQRSTUVWXYZ[\]^_`abcdefghijk  
!"#$%&'()*+,-./0123456789:;<=>?@ABCDEFGHIJKLMNOPQRSTUVWXYZ[\]^_`abcdefghijkl  
!"#$%&'()*+,-./0123456789:;<=>?@ABCDEFGHIJKLMNOPQRSTUVWXYZ[\]^_`abcdefghijklm  
:  
:
```

\r\n added

72 characters rotating out of 95 ASCII printable characters

SocketChannel

- ◉ java.nio.channels.SocketChannel
- ◉ Opening a channel

```
SocketAddress rama = new InetSocketAddress("rama.poly.edu", 19);  
SocketChannel client = SocketChannel.open(rama);
```

opened in blocking mode
(next line executed only after conn established)

- IOException is thrown if connection can't be established

Reading/Writing

◉ Reading

- input streams are not needed
- read directly from channel

```
ByteBuffer buffer = ByteBuffer.allocate(74);    why 74 bytes?  
int bytesRead = client.read(buffer);
```

read a sequence of bytes from client into buffer
read at least one byte or return -1 to indicate the end of data

◉ Writing

- output streams are not needed
- write directly to channel

```
WritableByteChannel output = Channels.newChannel(System.out);  
buffer.flip();    limit is set to current position; position is set to zero  
output.write(buffer);
```

don't need to tell output channel how many bytes to write; buffer keeps track of it

Reusing Buffer

- Creating buffer for every read/write operation can kill performance (slow down)
- Reuse existing buffer
 - need to reset position to zero
`buffer.clear();` doesn't delete data, but data will be overwritten
- flip() vs. clear()
 - flipping prepares for writing (write buffer data to channel)
 - clearing prepares for reading (write data read from channel to buffer)

```
import java.nio.*;
import java.nio.channels.*;
import java.net.*;
import java.io.IOException;

public class ChargenClient {

    public static int DEFAULT_PORT = 19;

    public static void main(String[] args) {

        if (args.length == 0) {
            System.out.println("Usage: java ChargenClient host [port]");
            return;
        }

        int port;
        try {
            port = Integer.parseInt(args[1]);
        } catch (RuntimeException ex) {
            port = DEFAULT_PORT;
        }

        try {
            SocketAddress address = new InetSocketAddress(args[0], port);
            SocketChannel client = SocketChannel.open(address);

            ByteBuffer buffer = ByteBuffer.allocate(74);
```

Example (contd.)

```
WritableByteChannel out = Channels.newChannel(System.out);

while (client.read(buffer) != -1) {
    buffer.flip();
    out.write(buffer);
    buffer.clear();
}
} catch (IOException ex) {
    ex.printStackTrace();
}
}
```

● Exercise

- save the buffer content to a file

Activating Nonblocking Mode

- The example code works in blocking mode
- Activating nonblocking mode

`client.configureBlocking(false);` don't block
true means block

- `read()` returns immediately even if there is no data available to read

- Modified code

```
while (true) {  
    // Put whatever code here you want to run every pass through the loop  
    // whether anything is read or not  
    int n = client.read(buffer);  
    if (n > 0) {  
        buffer.flip();  
        out.write(buffer);  
        buffer.clear();  
    } else if (n == -1) {  
        // This shouldn't happen unless the server is misbehaving.  
        break;  
    }  
}
```

need to check whether `read()` really read data or not

Example Server

Chargen Server

ServerSocketChannel

- Opening channel

```
ServerSocketChannel serverChannel = ServerSocketChannel.open();
```

this doesn't listen to any port (until it's bound to a port)

```
ServerSocket ss = serverChannel.socket();
```

```
ss.bind(new InetSocketAddress(19));
```

retrieve server socket and bind it to a port

alternatively

```
serverChannel.bind(new InetSocketAddress(19));
```

- Accepting request

```
SocketChannel clientChannel = serverChannel.accept();
```

this line will be placed differently from the case of blocking mode

Activating Nonblocking Mode

- Client channel

```
clientChannel.configureBlocking(false);
```

- for writing data to client
- write() returns when client channel is not ready to be written

- Server channel

```
serverChannel.configureBlocking(false);
```

- for accepting request from client
- accept() returns null when there is no request
- better to check if returned socket channel is null or not

Selector

- Enable program to iterate over all connections that are ready to be processed

```
Selector selector = Selector.open();
```

- Register each channel with selector (that monitors it)

```
serverChannel.register(selector, SelectionKey.OP_ACCEPT);
```

ready to accept new connection?

operation you are interested in

```
SelectionKey key = clientChannel.register(selector, SelectionKey.OP_WRITE);
```

ready to write?

each SelectionKey has an attachment of arbitrary Object type



Constructing Buffers

- In this example, we will attach buffer that channel writes onto network
- Once buffer is fully drained, we will refill it (reused)
- Buffer containing two sequence copies of data

```
byte[] rotation = new byte[95*2];  
for (byte i = ' '; i <= '~'; i++) {  
    rotation[i - ' '] = i;  
    rotation[i + 95 - ' '] = i;  
}
```

- this will be used to construct buffer containing 72 characters
- this form is convenient because what?

!"#\$%&'()*+,-./0123456789:;<=>?@ABCDEFGHIJKLMNOPQRSTUVWXYZ[\]^_`abcdefghijklmnopqrstuvwxyz{|}~ !"#\$%&'()*+,-./0123456789:;...

Constructing Buffers (contd.)

- Fill the buffer with 72 characters from rotation array and add `\r\n`

```
ByteBuffer buffer = ByteBuffer.allocate(74);  
buffer.put(rotation, 0, 72); write rotation[0~0+71] into buffer  
buffer.put((byte) '\r');  
buffer.put((byte) '\n');  
buffer.flip(); set position to zero  
key2.attach(buffer); attach this buffer to key
```

Checking Readiness

- Check whether anything is ready to be acted on

```
while (true) {  
    selector.select (); check  
    // process selected keys... get things that are ready  
    Set<SelectionKey> readyKeys = selector.selectedKeys();  
    Iterator iterator = readyKeys.iterator();  
    while (iterator.hasNext()) {  
        SelectionKey key = iterator.next();  
        // Remove key from set so we don't process it twice  
        iterator.remove();  
        // operate on the channel...  
    }  
}
```

- If ready channel is server channel, program accepts new socket channel and add it to selector
- If ready channel is socket channel, program writes as much of the buffer as it can onto the channel
- If no channels are ready, selector waits for one
- One thread(main thread) processes multiple simultaneous connections

Checking Operation Types

```
try {  
    if (key.isAcceptable()) {  
        ServerSocketChannel server = (ServerSocketChannel) key.channel();  
        SocketChannel connection = server.accept();  
        connection.configureBlocking(false);  
        connection.register(selector, SelectionKey.OP_WRITE);  
        // set up the buffer for the client...  
    } else if (key.isWritable()) {  
        SocketChannel client = (SocketChannel) key.channel();  
        // write data to client...  
    }  
}
```

Writing Data onto Channel

```
ByteBuffer buffer = (ByteBuffer) key.attachment(); get attachment
if (!buffer.hasRemaining()) { no elements between position and limit
    // Refill the buffer with the next line
    // Figure out where the last line started
    buffer.rewind(); move position to zero
    int first = buffer.get(); ready a byte (and increment position)
    // Increment to the next character
    buffer.rewind(); move position to zero
    int position = first - ' ' + 1; start from "!" +1 makes it rotating
    buffer.put(rotation, position, 72);
    buffer.put((byte) '\r');
    buffer.put((byte) '\n');
    buffer.flip();
}
client.write(buffer); this also updates buffer in the attachment of key
```

```
import java.nio.*;
import java.nio.channels.*;
import java.net.*;
import java.util.*;
import java.io.IOException;

public class ChargenServer {

    public static int DEFAULT_PORT = 19;

    public static void main(String[] args) {

        int port;
        try {
            port = Integer.parseInt(args[0]);
        } catch (RuntimeException ex) {
            port = DEFAULT_PORT;
        }
        System.out.println("Listening for connections on port " + port);

        byte[] rotation = new byte[95*2];
        for (byte i = ' '; i <= '~'; i++) {
            rotation[i - ' '] = i;
            rotation[i + 95 - ' '] = i;
        }
    }
}
```

```
ServerSocketChannel serverChannel;  
Selector selector;  
try {  
    serverChannel = ServerSocketChannel.open();  
    ServerSocket ss = serverChannel.socket();  
    InetSocketAddress address = new InetSocketAddress(port);  
    ss.bind(address);  
    serverChannel.configureBlocking(false);  
    selector = Selector.open();  
    serverChannel.register(selector, SelectionKey.OP_ACCEPT);  
} catch (IOException ex) {  
    ex.printStackTrace();  
    return;  
}  
  
while (true) {  
    try {  
        selector.select();  
    } catch (IOException ex) {  
        ex.printStackTrace();  
        break;  
    }  
}
```



```

Set<SelectionKey> readyKeys = selector.selectedKeys();
Iterator<SelectionKey> iterator = readyKeys.iterator();
while (iterator.hasNext()) {

    SelectionKey key = iterator.next();
    iterator.remove();
    try {
        if (key.isAcceptable()) {
            ServerSocketChannel server = (ServerSocketChannel) key.channel();
            SocketChannel client = server.accept();
            System.out.println("Accepted connection from " + client);
            client.configureBlocking(false);
            SelectionKey key2 = client.register(selector, SelectionKey.
                                                         OP_WRITE);
            ByteBuffer buffer = ByteBuffer.allocate(74);
            buffer.put(rotation, 0, 72);
            buffer.put((byte) '\r');
            buffer.put((byte) '\n');
            buffer.flip();
            key2.attach(buffer);
        }
    }
}

```

```

} else if (key.isWritable()) {
    SocketChannel client = (SocketChannel) key.channel();
    ByteBuffer buffer = (ByteBuffer) key.attachment();
    if (!buffer.hasRemaining()) {
        // Refill the buffer with the next line
        buffer.rewind();
        // Get the old first character
        int first = buffer.get();
        // Get ready to change the data in the buffer
        buffer.rewind();
        // Find the new first characters position in rotation
        int position = first - ' ' + 1;
        // copy the data from rotation into the buffer
        buffer.put(rotation, position, 72);
        // Store a line break at the end of the buffer
        buffer.put((byte) '\r');
        buffer.put((byte) '\n');
        // Prepare the buffer for writing
        buffer.flip();
    }
    client.write(buffer);
}

```

```
    } catch (IOException ex) {
        key.cancel();
        try {
            key.channel().close();
        }
        catch (IOException cex) {}
    }
}
}
```

Notes

- ◉ Can extend this to the case of multiple threads
 - multiple CPUs can be exploited
- ◉ Can also extend to the case of thread pool
- ◉ `select()` ensures you are never wasting any time on connections that are not ready to receive data

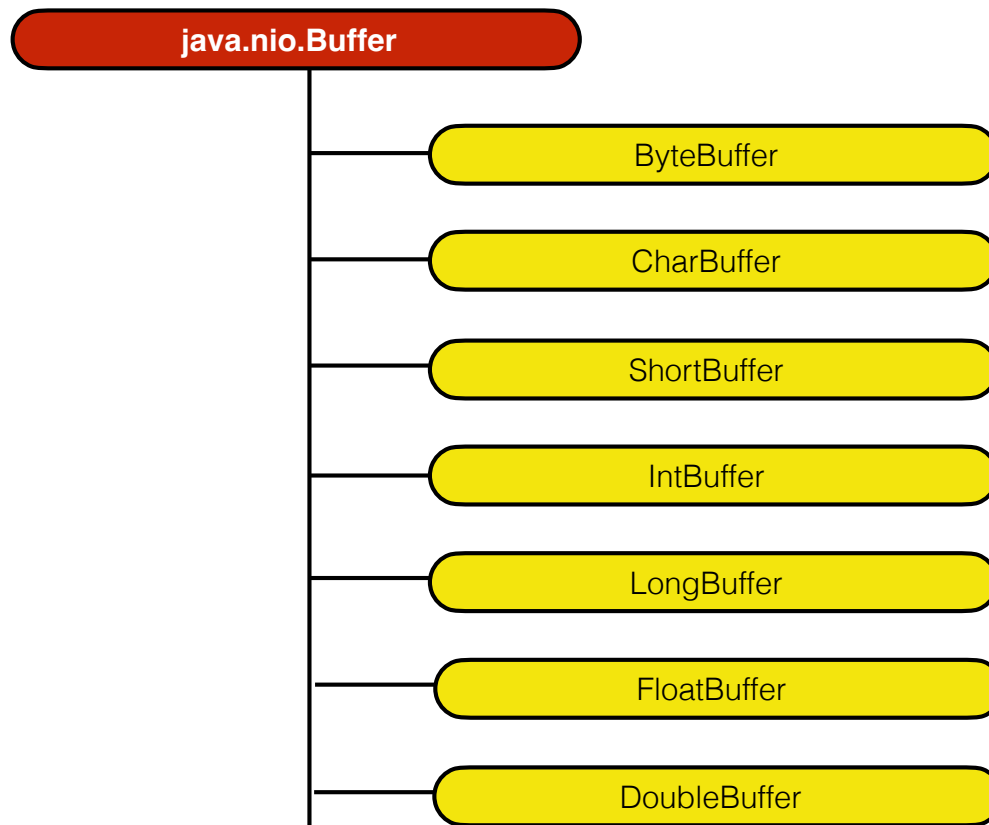
Buffers

Streams vs. Channels

- ◉ Streams are byte-based
 - data is written or read in bytes
- ◉ Channels are block-based
 - data is written or read one buffer at a time
- ◉ Streams
 - read or write only
- ◉ Channels
 - can both read and write

Class Hierarchy

- Container for data of specific type



Network programs use `ByteBuffer` almost exclusively

Four Properties

◎ position

- next location in buffer that will be read from or written to
- start counting at 0, and max value=size of buffer
- methods

```
public final int    position()  
public final Buffer position(int newPosition)
```

◎ capacity

- max number of elements buffer can hold
- set when buffer is created
- method

```
public final int capacity()
```


Four Properties (contd.)

◎ limit

- end of accessible data
- cannot write or read at or past this point

```
public final int    limit()  
public final Buffer limit(int newLimit)
```

◎ mark


- client-specified index

```
public final Buffer mark()    mark is set at current position  
public final Buffer reset()  current position is set to marked position
```

- if position is set below an existing mark, mark is discarded

Initialization

- Position is set to zero and mark is undefined
- Each element of a newly-created buffer is initialized to zero
- Invariants
 -
- Exercise
 - create a ByteBuffer object with capacity 10, and check its initial properties





Clearing

- `clear()`

```
public final Buffer clear()
```

```
ByteBuffer buf = ByteBuffer.allocate(10);  
...  
buf.clear();
```





- get ready for being written

Rewinding

- `rewind()`

```
public final Buffer rewind()
```

```
ByteBuffer buf = ByteBuffer.allocate(10);  
...  
buf.rewind();
```



- allow the buffer to be reread


Flipping

• flip()


```
public final Buffer flip()
```

```
ByteBuffer buf = ByteBuffer.allocate(10);  
...  
buf.flip();
```

Before flip()



After flip()



- get ready for reading from the buffer just filled

Info about Buffer

- Remaining elements

```
public final int    remaining()
```

number of elements between position and limit

```
public final boolean hasRemaining()
```

true if remaining()>0

Creating Buffers: Allocation

● Allocation

- allocate() returns a new, empty buffer with fixed specified capacity

```
ByteBuffer buffer1 = ByteBuffer.allocate(100);  
IntBuffer  buffer2 = IntBuffer.allocate(100);
```

- implemented on top of Java array and can be accessed by array() and arrayOffset() ——— return offset in the backing array

```
byte[] data1 = buffer1.array();  
int[]  data2 = buffer2.array();
```

- changes in data1, data2 are reflected in buffer1 and buffer2

● Exercise

- check whether changes in int array can be seen in IntBuffer

Creating Buffer: Wrapping

- In case you already have an array of data, you just need to wrap it to create a buffer

```
byte[] data = "Some data".getBytes("UTF-8");  
ByteBuffer buffer1 = ByteBuffer.wrap(data);  
char[] text = "Some text".toArray();  
CharBuffer buffer2 = CharBuffer.wrap(text);
```


Creating Buffers: Direct Allocation

- ByteBuffer class has an additional method `allocateDirect()`

- don't create a backing array

```
ByteBuffer buffer = ByteBuffer.allocateDirect(100);
```

- direct memory access to the buffer on an Ethernet card, kernel memory, or something else

- Exercise

- try to call `array()` and `arrayOffset()`



Filling and Draining

- based on current position
 - public abstract byte get()
 - public abstract ByteBuffer put(byte b)
- based on specified position
 - public abstract byte get(int index)
 - public abstract ByteBuffer put(int index, byte b)
- using array
 - public ByteBuffer get(byte[] dst)
 - public ByteBuffer get(byte[] dst, int offset, int length)
 - public final ByteBuffer put(byte[] src)
 - public ByteBuffer put (byte[] src, int offset, int length)
- buffer as an input argument
 - public ByteBuffer put (ByteBuffer src)

get() method

- read 1 byte and increment position


```
ByteBuffer buf = ByteBuffer.allocate(10);  
...  
byte b = buf.get();
```



put() method

- write 1 byte and increment position

```
ByteBuffer buf = ByteBuffer.allocate(10);  
...  
byte b = buf.put((byte) z);
```



- Exercise

- what's the position?

```
CharBuffer buffer = CharBuffer.allocate(12);  
buffer.put('H');  
buffer.put('e');  
buffer.put('l');  
buffer.put('l');  
buffer.put('o');
```

Example

```
ChargenServer.ja  ChargenServer.j  ChargenClient.j  ChargenClientEx  *BufferProperty

import java.nio.ByteBuffer;

public class RelativeBufferTest {

    public static void main(String[] args) {
        // TODO Auto-generated method stub
        ByteBuffer buf = ByteBuffer.allocate(10);

        System.out.print("Init Position: " + buf.position());
        System.out.print(", Init Limit: " + buf.limit());
        System.out.println(", Init Capacity: " + buf.capacity());

        buf.mark();

        buf.put((byte) 10).put((byte) 11).put((byte) 12);


        buf.reset();

        System.out.println("Value: " + buf.get() + ", Position: " + buf.position());
        System.out.println("Value: " + buf.get() + ", Position: " + buf.position());
        System.out.println("Value: " + buf.get() + ", Position: " + buf.position());
        System.out.println("Value: " + buf.get() + ", Position: " + buf.position());
    }
}
```


get(int index)

```
ByteBuffer buf = ByteBuffer.allocate(10);  
...  
byte b = buf.get(5);
```

Before get()




After get()




put(int index, byte b)

```
ByteBuffer buf = ByteBuffer.allocate(10);  
...  
byte b = buf.put(5, (byte) z);
```

Before put()



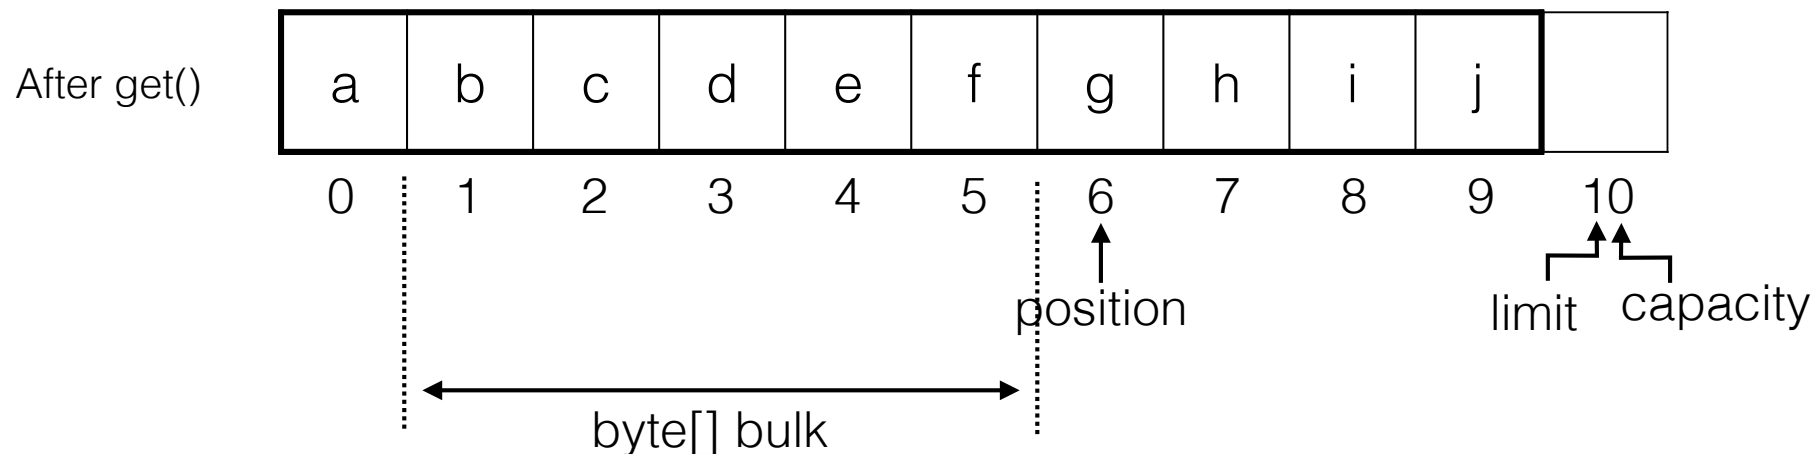

After put()



Bulk Methods

```
ByteBuffer buf = ByteBuffer.allocate(10);  
byte[] bulk = new byte[5];  
...  
byte b = buf.get(bulk);
```

`get(bulk) == get(bulk, 0, bulk.length)`



Example

```
import java.nio.ByteBuffer;

public class BulkReadTest {

    public static void main(String[] args) {
        // TODO Auto-generated method stub
        ByteBuffer buf = ByteBuffer.allocate(10);
        buf.put((byte) 0).put((byte) 1).put((byte) 2).put((byte) 3).put((byte) 4);
        buf.mark();
        buf.put((byte) 5).put((byte) 6).put((byte) 7).put((byte) 8).put((byte) 9);
        buf.reset();

        byte[] b = new byte[15];

        int size = buf.remaining();
        if (b.length < size) {
            size = b.length;
        }

        System.out.println("Position: " + buf.position() + ", Limit: " + buf.limit());

        buf.get(b, 0, size);

        System.out.println("Position: " + buf.position() + ", Limit: " + buf.limit());

        doSomething(b, size);
    }

    public static void doSomething(byte[] b, int size) {
        for (int i = 0; i < size; i++) {
            System.out.println("byte = " + b[i]);
        }
    }
}
```

● Exercise

- try to read more bytes than remaining bytes in buf

Example

```
import java.nio.ByteBuffer;

public class BulkWriteTest {

    public static void main(String[] args) {
        // TODO Auto-generated method stub
        ByteBuffer buf = ByteBuffer.allocate(10);
        buf.position(5);
        buf.mark();
        System.out.println("Position: " + buf.position() + ", Limit: " + buf.limit());

        byte[] b = new byte[15];
        for (int i = 0; i < b.length; i++) {
            b[i] = (byte) i;
        }

        int size = buf.remaining();
        if (b.length < size) {
            size = b.length;
        }

        System.out.println("Position: " + buf.position() + ", Limit: " + buf.limit());

        buf.put(b, 0, size);

        System.out.println("Position: " + buf.position() + ", Limit: " + buf.limit());

        doSomething(buf, size);
    }


    public static void doSomething(ByteBuffer buf, int size) {
        for (int i = 0; i < size; i++) {
            System.out.println("byte = " + buf.get());
        }
    }
}
```

● Exercise

- fix the BufferUnderflowException

Data Conversion

- All data in Java ultimately resolves to bytes
 - any primitive data type (like int, double, float,...) can be written as bytes
- Any sequence of bytes of right length can be interpreted as a primitive datum



Methods

<code>public abstract char</code>	<code>getChar()</code>	- read next two bytes and compose them into a character
<code>public abstract ByteBuffer</code>	<code>putChar(char value)</code>	- increment position by two
<code>public abstract char</code>	<code>getChar(int index)</code>	- write two bytes into current position
<code>public abstract ByteBuffer</code>	<code>putChar(int index, char value)</code>	- increment position by two
<code>public abstract short</code>	<code>getShort()</code>	
<code>public abstract ByteBuffer</code>	<code>putShort(short value)</code>	
<code>public abstract short</code>	<code>getShort(int index)</code>	
<code>public abstract ByteBuffer</code>	<code>putShort(int index, short value)</code>	
<code>public abstract int</code>	<code>getInt()</code>	
<code>public abstract ByteBuffer</code>	<code>putInt(int value)</code>	
<code>public abstract int</code>	<code>getInt(int index)</code>	
<code>public abstract ByteBuffer</code>	<code>putInt(int index, int value)</code>	
<code>public abstract long</code>	<code>getLong()</code>	
<code>public abstract ByteBuffer</code>	<code>putLong(long value)</code>	
<code>public abstract long</code>	<code>getLong(int index)</code>	
<code>public abstract ByteBuffer</code>	<code>putLong(int index, long value)</code>	
<code>public abstract float</code>	<code>getFloat()</code>	
<code>public abstract ByteBuffer</code>	<code>putFloat(float value)</code>	
<code>public abstract float</code>	<code>getFloat(int index)</code>	
<code>public abstract ByteBuffer</code>	<code>putFloat(int index, float value)</code>	
<code>public abstract double</code>	<code>getDouble()</code>	
<code>public abstract ByteBuffer</code>	<code>putDouble(double value)</code>	
<code>public abstract double</code>	<code>getDouble(int index)</code>	
<code>public abstract ByteBuffer</code>	<code>putDouble(int index, double value)</code>	

Example

```
import java.nio.*;

public class DataConversionTest {

    public static void main(String[] args) {
        // TODO Auto-generated method stub
        ByteBuffer buf = ByteBuffer.allocate(16);
        int i = 0;
        while (buf.hasRemaining()) {
            buf.put((byte)(i));
            i++;
        }
        buf.flip();

        System.out.println(buf);
        showBuffer(buf, "int");
        showBuffer(buf, "char");
        showBuffer(buf, "float");
        showBuffer(buf, "long");
    }

    public static void showBuffer(ByteBuffer buf, String type) {
        if (type.equals("int")) {
            while (buf.hasRemaining()) {
                System.out.println(buf.getInt());
            }
            buf.flip();
        } else if (type.equals("char")) {
            while (buf.hasRemaining()) {
                System.out.println(buf.getChar());
            }
            buf.flip();
        } else if (type.equals("float")) {
            while (buf.hasRemaining()) {
                System.out.println(buf.getFloat());
            }
            buf.flip();
        } else if (type.equals("long")) {
            while (buf.hasRemaining()) {
                System.out.println(buf.getLong());
            }
            buf.flip();
        }
    }
}
```

Example: IntgenServer

```
import java.nio.*;
import java.nio.channels.*;
import java.net.*;
import java.util.*;
import java.io.IOException;

public class IntgenServer {

    public static int DEFAULT_PORT = 1919;

    public static void main(String[] args) {

        int port;
        try {
            port = Integer.parseInt(args[0]);
        } catch (RuntimeException ex) {
            port = DEFAULT_PORT;
        }
        System.out.println("Listening for connections on port " + port);

        ServerSocketChannel serverChannel;
        Selector selector;
        try {
            serverChannel = ServerSocketChannel.open();
            ServerSocket ss = serverChannel.socket();
            InetSocketAddress address = new InetSocketAddress(port);
            ss.bind(address);
            serverChannel.configureBlocking(false);
            selector = Selector.open();
            serverChannel.register(selector, SelectionKey.OP_ACCEPT);
        } catch (IOException ex) {
            ex.printStackTrace();
            return;
        }
    }
}
```

```

while (true) {
    try {
        selector.select();
    } catch (IOException ex) {
        ex.printStackTrace();
        break;
    }

    Set<SelectionKey> readyKeys = selector.selectedKeys();
    Iterator<SelectionKey> iterator = readyKeys.iterator();
    while (iterator.hasNext()) {
        SelectionKey key = iterator.next();
        iterator.remove();
        try {
            if (key.isAcceptable()) {
                ServerSocketChannel server = (ServerSocketChannel) key.channel();
                SocketChannel client = server.accept();
                System.out.println("Accepted connection from " + client);
                client.configureBlocking(false);
                SelectionKey key2 = client.register(selector, SelectionKey.OP_WRITE);
                ByteBuffer output = ByteBuffer.allocate(4);
                output.putInt(0);
                output.flip();
                key2.attach(output);
            } else if (key.isWritable()) {
                SocketChannel client = (SocketChannel) key.channel();
                ByteBuffer output = (ByteBuffer) key.attachment();
                if (!output.hasRemaining()) {
                    output.rewind();
                    int value = output.getInt();
                    output.clear();
                    output.putInt(value + 1);
                    output.flip();
                }
                client.write(output);
            }
        } catch (IOException ex) {
            key.cancel();
            try {
                key.channel().close();
            }
            catch (IOException cex) {}
        }
    }
}
}
}

```

View Buffers

- Interface for viewing buffer in a particular data type

- Methods

```
public abstract ShortBuffer asShortBuffer()  
public abstract CharBuffer  asCharBuffer()  
public abstract IntBuffer   asIntBuffer() create a view of ByteBuffer as an IntBuffer  
public abstract LongBuffer  asLongBuffer()  
public abstract FloatBuffer asFloatBuffer()  
public abstract DoubleBuffer asDoubleBuffer()
```

- each buffer has its own limit, capacity, mark and position
- changes in view are reflected in underlying buffer and vice versa

- Example

```
ByteBuffer buffer = ByteBuffer.allocate(4);  
IntBuffer  view  = buffer.asIntBuffer();
```


Example: IntgenClient

```
import java.nio.*;
import java.nio.channels.*;
import java.net.*;
import java.io.IOException;

public class IntgenClient {

    public static int DEFAULT_PORT = 1919;

    public static void main(String[] args) {

        if (args.length == 0) {
            System.out.println("Usage: java IntgenClient host [port]");
            return;
        }

        int port;
        try {
            port = Integer.parseInt(args[1]);
        } catch (RuntimeException ex) {
            port = DEFAULT_PORT;
        }

        try {
            SocketAddress address = new InetSocketAddress(args[0], port);
            SocketChannel client = SocketChannel.open(address);
            ByteBuffer buffer = ByteBuffer.allocate(4);
            IntBuffer view = buffer.asIntBuffer();

            for (int expected = 0; ; expected++) {
                client.read(buffer);
                int actual = view.get();
                buffer.clear();
                view.rewind();

                if (actual != expected) {
                    System.err.println("Expected " + expected + "; was " + actual);
                    break;
                }
                System.out.println(actual);
            }
        } catch (IOException ex) {
            ex.printStackTrace();
        }
    }
}
```

Compacting Buffers


- Methods

```
public abstract ByteBuffer compact()  
public abstract IntBuffer compact()  
public abstract ShortBuffer compact()  
public abstract FloatBuffer compact()  
public abstract CharBuffer compact()  
public abstract DoubleBuffer compact()
```

useful when write is not finished
and read must begin

Compacting Buffers

```
ByteBuffer buf = ByteBuffer.allocate(10);  
...  
buf.compact();
```




Duplicating Buffers

- Create a copy of a buffer to deliver the same information
 - same backing array
 - change in one buffer is reflected to all other copies and backing array (so mainly used for reading)
- Useful when transmitting same data over multiple channels

Duplicating Buffers

```
ByteBuffer buf = ByteBuffer.allocate(10);  
...  
ByteBuffer buf2 = buf.duplicate();
```



Example

- ◉ Nonblocking version SingleFileHTTPServer
- ◉ Single file is stored in one constant, read-only buffer
- ◉ Every time a client connects, the program makes a duplicate of this buffer just for that channel (which is stored as channel's attachment)
 - without duplicates, one client has to wait until the other finishes so the original buffer can be rewound
 - duplicates enable simultaneous buffer reuse

```

import java.io.*;
import java.nio.*;
import java.nio.channels.*;
import java.nio.charset.*;
import java.nio.file.*;
import java.util.*;
import java.net.*;

public class NonblockingSingleFileHTTPServer {

    private ByteBuffer contentBuffer;
    private int port = 80;

    public NonblockingSingleFileHTTPServer(
        ByteBuffer data, String encoding, String MIMEType, int port) {

        this.port = port;
        String header = "HTTP/1.0 200 OK\r\n"
            + "Server: NonblockingSingleFileHTTPServer\r\n"
            + "Content-length: " + data.limit() + "\r\n"
            + "Content-type: " + MIMEType + "\r\n\r\n";
        byte[] headerData = header.getBytes(Charset.forName("US-ASCII"));

        ByteBuffer buffer = ByteBuffer.allocate(
            data.limit() + headerData.length);
        buffer.put(headerData);
        buffer.put(data);
        buffer.flip();
        this.contentBuffer = buffer;
    }

    public void run() throws IOException {
        ServerSocketChannel serverChannel = ServerSocketChannel.open();
        ServerSocket serverSocket = serverChannel.socket();
        Selector selector = Selector.open();
        InetSocketAddress localPort = new InetSocketAddress(port);
        serverSocket.bind(localPort);
        serverChannel.configureBlocking(false);
        serverChannel.register(selector, SelectionKey.OP_ACCEPT);
    }
}

```

```

while (true) {
    selector.select();
    Iterator<SelectionKey> keys = selector.selectedKeys().iterator();
    while (keys.hasNext()) {
        SelectionKey key = keys.next();
        keys.remove();
        try {
            if (key.isAcceptable()) {
                ServerSocketChannel server = (ServerSocketChannel) key.channel();
                SocketChannel channel = server.accept();
                channel.configureBlocking(false);
                channel.register(selector, SelectionKey.OP_READ);
            } else if (key.isWritable()) {
                SocketChannel channel = (SocketChannel) key.channel();
                ByteBuffer buffer = (ByteBuffer) key.attachment();
                if (buffer.hasRemaining()) {
                    channel.write(buffer);
                } else { // we're done
                    channel.close();
                }
            } else if (key.isReadable()) {
                // Don't bother trying to parse the HTTP header.
                // Just read something.
                SocketChannel channel = (SocketChannel) key.channel();
                ByteBuffer buffer = ByteBuffer.allocate(4096);
                channel.read(buffer);
                // switch channel to write-only mode
                key.interestOps(SelectionKey.OP_WRITE);
                key.attach(contentBuffer.duplicate());
            }
        } catch (IOException ex) {
            key.cancel();
            try {
                key.channel().close();
            }
            catch (IOException cex) {}
        }
    }
}

```



```

public static void main(String[] args) {
    if (args.length == 0) {
        System.out.println(
            "Usage: java NonblockingSingleFileHTTPServer file port encoding");
        return;
    }

    try {
        // read the single file to serve
        String contentType =
            URLConnection.getFileNameMap().getContentTypeFor(args[0]);
        Path file = FileSystems.getDefault().getPath(args[0]);
        byte[] data = Files.readAllBytes(file);
        ByteBuffer input = ByteBuffer.wrap(data);

        // set the port to listen on
        int port;
        try {
            port = Integer.parseInt(args[1]);
            if (port < 1 || port > 65535) port = 80;
        } catch (RuntimeException ex) {
            port = 80;
        }

        String encoding = "UTF-8";
        if (args.length > 2) encoding = args[2];

        NonblockingSingleFileHTTPServer server
            = new NonblockingSingleFileHTTPServer(
                input, encoding, contentType, port);
        server.run();
    } catch (IOException ex) {
        System.err.println(ex);
    }
}
}

```

Exercise: modify the code for korean text

Read-only Copy

- `asReadOnlyBuffer()`

- same as `duplicate()` except that copy created by `asReadOnlyBuffer()` only allows reading


- `slice()`

- copy a subsequence of original buffer from position to limit
- have independent 4 properties

```
public abstract ByteBuffer slice()  
public abstract IntBuffer slice()  
public abstract ShortBuffer slice()  
public abstract FloatBuffer slice()  
public abstract CharBuffer slice()  
public abstract DoubleBuffer slice()
```

Read-only Copy

```
ByteBuffer buf = ByteBuffer.allocate(10);  
...  
ByteBuffer buf2 = buf.slice();
```



Channels

Channels

- Channels move blocks of data into and out of buffers to and from various I/O sources like files, sockets, datagrams,...
- Three main channel classes
 - SocketChannel: TCP
 - ServerSocketChannel: TCP
 - DatagramChannel: UDP

SocketChannel

SocketChannel Class

- ◉ Read from and write to TCP sockets
- ◉ Data must be encoded in ByteBuffer
- ◉ Connecting

```
public static SocketChannel open(SocketAddress remote) throws IOException  
public static SocketChannel open() throws IOException
```

- blocking mode example

```
SocketAddress address = new InetSocketAddress("www.cafeaulait.org", 80);  
SocketChannel channel = SocketChannel.open(address); open connection
```

- open(address) blocks

- another example

```
SocketChannel channel = SocketChannel.open(); create unconnected socket  
SocketAddress address = new InetSocketAddress("www.cafeaulait.org", 80);  
channel.connect(address);
```

SocketChannel Class (contd.)

- Nonblocking mode example

```
SocketChannel channel = SocketChannel.open();  
SocketAddress address = new InetSocketAddress("www.cafeaulait.org", 80);  
channel.configureBlocking(false);  
channel.connect(); return immediately even before connection is established
```

- Other methods

```
public abstract boolean isConnected()  
public abstract boolean isConnectionPending()
```

- isConnected(): true if connection established (socket is open and connected)
- isConnectionPending(): true if connection is still being set up, but not yet open

Reading

● Method

```
public abstract int read(ByteBuffer dst) throws IOException
```

- channel fills buffer with as much data as it can
- return # bytes put in the buffer
- when it encounters end of stream while reading, it returns -1 on the next call to read()
- blocking
 - read at least one byte, return -1, or throw an exception
- nonblocking
 - can return 0

Reading (contd.)

- How to read until buffer is filled or end of stream is detected?

- Scattering

```
public final long read(ByteBuffer[] dsts) throws IOException  
public final long read(ByteBuffer[] dsts, int offset, int length)  
    throws IOException
```

- fill several buffers from one source
- example

```
ByteBuffer[] buffers = new ByteBuffer[2];  
buffers[0] = ByteBuffer.allocate(1000);  
buffers[1] = ByteBuffer.allocate(1000);  
while (buffers[1].hasRemaining() && channel.read(buffers) != -1) ;
```

Writing

- Method

```
public abstract int write(ByteBuffer src) throws IOException
```

- doesn't guarantee to write complete content of buffer in nonblocking mode
- useful clause

```
while (buffer.hasRemaining() && channel.write(buffer) != -1) ;
```

- Buffer array

```
public final long write(ByteBuffer[] srcs) throws IOException  
public final long write(ByteBuffer[] srcs, int offset, int length)  
    throws IOException
```

offset in ByteBuffer array index

- for example, HTTP head and body in separate buffers

Closing

- Methods

```
public void close() throws IOException
```

```
public boolean isOpen()
```

- true if channel is open

ServerSocketChannel Class

ServerSocketChannel

- One purpose
 - accept incoming connections
- This class declares only four methods including accept()
- Creating

```
try {  
    • ServerSocketChannel server = ServerSocketChannel.open(); just create object  
      ServerSocket socket = serverChannel.socket();  
      SocketAddress address = new InetSocketAddress(80);  
      socket.bind(address);  
} catch (IOException ex) {  
    System.err.println("Could not bind to port 80 because " + ex.getMessage());  
}
```

or

```
try {  
    ServerSocketChannel server = ServerSocketChannel.open();  
    SocketAddress address = new InetSocketAddress(80);  
    server.bind(address);  
} catch (IOException ex) {  
    System.err.println("Could not bind to port 80 because " + ex.getMessage());  
}
```

Accepting

- Method

```
public abstract SocketChannel accept() throws IOException
```

```
ServerSocketChannel server = (ServerSocketChannel) key.channel();  
SocketChannel client = server.accept();
```

- Blocking mode (default)

- wait for incoming connection, and
- return SocketChannel object connected to client

- Nonblocking mode

- return null if no incoming connection
- appropriate for servers doing a lot of work for each connection and thus may want to process multiple requests in parallel
- need to call `configureBlocking(false)`

Channels Class

Channels Class

- Simple utility class for wrapping channels around traditional I/O-based streams, readers, writers,...
- Useful when you want to use the new I/O model in one part, but still incorporate legacy APIs that expect streams
- Methods

```
public static InputStream newInputStream(ReadableByteChannel ch)
public static OutputStream newOutputStream(WritableByteChannel ch)
public static ReadableByteChannel newChannel(InputStream in)
public static WritableByteChannel newChannel(OutputStream out)
public static Reader newReader (ReadableByteChannel channel,
    CharsetDecoder decoder, int minimumBufferCapacity)
public static Reader newReader (ReadableByteChannel ch, String encoding)
public static Writer newWriter (WritableByteChannel ch, String encoding)
```

```
WritableByteChannel out = Channels.newChannel(System.out);
```

Readiness Selection

Motivation

- It's important to choose a socket that will not block when read or written
 - primarily of interest to servers
- In order to perform readiness selection, channels are registered with a Selector object
 - each channel is assigned a SelectionKey
- Program can ask the Selector object for the set of keys to the channels that are ready to perform the operation you want to perform without blocking

Selector Class

- Creating a new selector, by calling `Selector.open()`

```
public static Selector open() throws IOException
```

```
Selector selector = Selector.open();
```

- Adding a channel to the selector

```
public final SelectionKey register(Selector sel, int ops)
    throws ClosedChannelException
public final SelectionKey register(Selector sel, int ops, Object att)
    throws ClosedChannelException
```

type of operation

selector channel is registering with

attachment

```
serverChannel.register(selector, SelectionKey.OP_ACCEPT);
SelectionKey key2 = client.register(selector, SelectionKey.OP_WRITE);
```

- defined in `SelectableChannel` class
 - not all channels are selectable(e.g., `FileChannel`) but all network channels are

- Types of operation (defined in `SelectionKey`)

```
SelectionKey.OP_ACCEPT
```

```
SelectionKey.OP_CONNECT
```

```
SelectionKey.OP_READ
```

```
SelectionKey.OP_WRITE
```

bit-flag int constants (1, 2, 4,...)

exercise: find out the number of each type

● Example

```
channel.register(selector, SelectionKey.OP_READ | SelectionKey.OP_WRITE);
```

- | bitwise OR operator
- register for both reading and writing on a socket

Selecting Ready Channels

- `selectNow()`

```
public abstract int selectNow() throws IOException
```

- perform a nonblocking select, i.e., immediately return if no connections are ready to be processed

- `select()`

```
public abstract int select() throws IOException      selector.select();  
public abstract int select(long timeout) throws IOException
```

timeout in ms => return 0 if no channels are ready in timeout

- wait until at least one registered channel is ready to be processed

- Retrieving ready channels

```
public abstract Set<SelectionKey> selectedKeys()
```

```
Set<SelectionKey> readyKeys = selector.selectedKeys();
```

- you need to iterate through the returned set

Closing

- Method

```
public abstract void close() throws IOException
```

- release any resources associated with selector
- cancel all keys registered
- interrupt up any threads blocked by one of this selector's select methods

SelectionKey Class

- Serve as pointers to channels
- Hold an object attachment
- SelectionKey objects are returned by register() method
 - single channel can be registered with multiple selectors
- Testing what a key is ready to do

```
public final boolean isAcceptable()  
public final boolean isConnectable()  
public final boolean isReadable()  
public final boolean isWritable()
```

- Retrieving a channel and attachment

```
• public abstract SelectableChannel channel()  
  public final Object attachment()      ServerSocketChannel server = (ServerSocketChannel) key.channel();
```


Canceling

◎ Method

```
public abstract void cancel()
```

- when you are finished with a connection, deregister its SelectionKey object so the selector doesn't waste any resources querying it for readiness
- not absolutely necessary

◎ Note

- closing a channel automatically deregisters all keys for that channel in all selectors
 - * a single channel can be registered with multiple selectors